

Ravani Roshan

AI Engineer · LLM Systems Developer · Agent Infrastructure Builder
Ahmedabad, Gujarat, India · +91 91060 95375 · ravaniroshansingh@gmail.com · github.com/RavaniRoshan

PROFESSIONAL SUMMARY

Computer Science Engineering undergraduate focused on AI engineering, LLM systems, agent infrastructure, RAG pipelines, and fine-tuning workflows. Built open-source tools including a hermetic multi-agent coding system, an in-process AI SDK guardrail layer, and a Windows-native MCP automation server for AI agents. Strong hands-on experience with Python, Rust, FastAPI, LangChain, LangGraph, MCP, Docker, PostgreSQL, and cloud deployment.

EDUCATION

Silver Oak University Ahmedabad, Gujarat, India
B.Tech, Computer Science Engineering CGPA: 8.0

EXPERIENCE

Google Developers Group Oct 2023 – Dec 2025
Core Developer Ahmedabad, Gujarat, India

- Led AI/tooling sessions for a 200+ developer community covering LLM integration, agent workflows, deployment patterns, and AI developer tools.
- Mentored peers on building practical AI applications with LangChain, LangGraph, APIs, and production deployment.
- Contributed to community projects across AI tooling, applied ML, and full-stack systems.

PROJECTS

Niki – Hermetic Multi-Agent Coding System | *Rust, Docker, CLI, Multi-agent, LLM* Jul 2026 – Present

- Built an autonomous multi-agent coding framework: agents plan, implement, test, and review code in isolated Docker sandboxes without human intervention.
- Implemented fire-and-forget agent orchestration with bring-your-own-key LLM routing, structured review loops, and git-branch-per-task output.

backstop – In-Process AI SDK Guardrails | *Python, OpenAI/Anthropic SDK, circuit-breaking* Jul 2026 – Present

- Built an in-process AI SDK safety layer enforcing token budgets, circuit-breaking, retry policies, and real-time metrics for OpenAI/Anthropic clients via transport-level interception.
- Implemented AIMD congestion control, priority-based admission, gzip-cached replay, and a multi-agent diff CLI (Wedge) for parallel agent cost isolation without proxy hops.

WinScript MCP – Windows Automation Server for AI Agents | *Python, MCP, FastMCP, SQLite* Apr 2026 – Present

- Built an MCP server exposing 59 tools for AI agents to control Windows apps, files, clipboard, screenshots, Office automation, and system state.
- Integrated UI Automation, COM, Win32, OCR fallback, SQLite audit logs, workflow replay, and safe/standard execution modes.
- Designed state-aware tool responses returning action results, window changes, confidence scores, and failure context.

Mnemo – Cognitive Telemetry for LLM Agents | *Python, LLM Agents, Observability* May 2026 – Present

- Built telemetry infrastructure for memory-augmented agents with trace recording, failure attribution, belief tracking, memory replay, and metrics reporting.
- Implemented recency-weighted attribution, memory dependency linking, contradiction monitoring, and replay/ablation analysis.

TECHNICAL SKILLS

AI/LLM: LLM applications, agent systems, LangChain, LangGraph, MCP, RAG, prompt orchestration, tool use, QLoRA, DPO, AI observability
Backend/Systems: Python, Rust, FastAPI, Node.js, Express, REST APIs, WebSockets, Docker, PostgreSQL, MySQL, Supabase, SQLite, Redis/Vercel KV
Frontend/Product: JavaScript, TypeScript, React, Next.js, Tailwind CSS, shadcn/ui, dashboards, documentation sites
Cloud/Tools: Vercel, Railway, Fly.io, AWS, GCP, GitHub Actions, npm, PyPI, Linux/Windows automation

CERTIFICATIONS

Google: MLOps for Generative AI (Dec 2025), Responsible AI: Privacy & Safety (Jan 2026), Transformer Models and BERT Model (Mar 2026)

Anthropic: Claude Code in Action (Mar 2026)

DeepLearning.AI: AI Agents in LangGraph, Finetuning Large Language Models, LangChain for LLM Application Development

Other: Machine Learning with Python (freeCodeCamp / Cognitive Class, Dec 2024)